

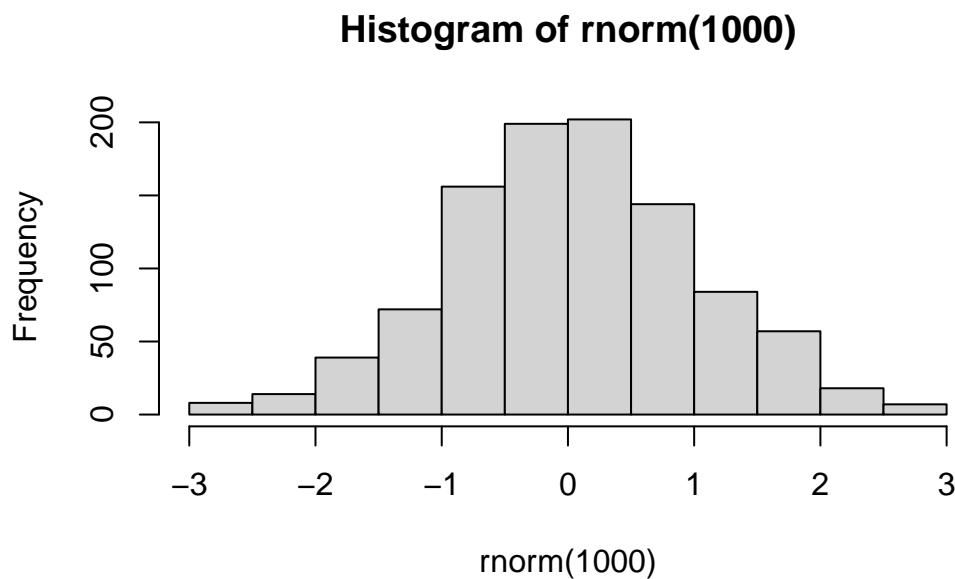
# class07: Machine Learning 1

Rachel (PID:A16037641)

Today we will begin our exploration of some “classical” machine learning approaches. We will start with clustering:

Let’s first make up some data to cluster where we know what the answer should be.

```
hist(rnorm(1000))
```

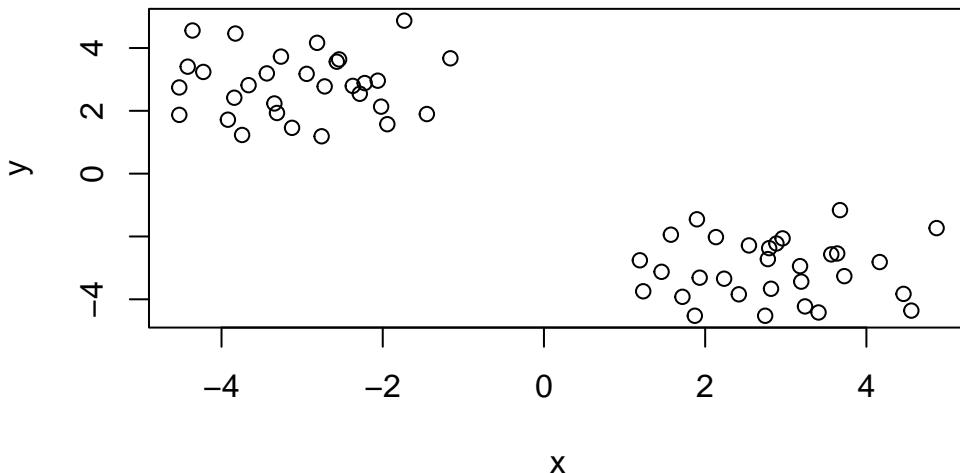


```
x <- c( rnorm(30, mean=-3), rnorm(30, mean=+3))
y<- rev(x)
x <- cbind(x,y)
head(x)
```

```
x      y
[1,] -3.311719 1.931786
[2,] -4.419169 3.406210
[3,] -2.814196 4.164933
[4,] -3.262510 3.726027
[5,] -2.720162 2.777510
[6,] -4.359768 4.557326
```

Peek at x with `plot()`

```
plot(x)
```



Then main function in “base” R for K-means clustering is called `kmeans()`.

```
k <- kmeans(x, centers=2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```
x      y
1  2.828172 -3.037670
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"       "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. How big are the clusters (i.e their size)?

k\$size

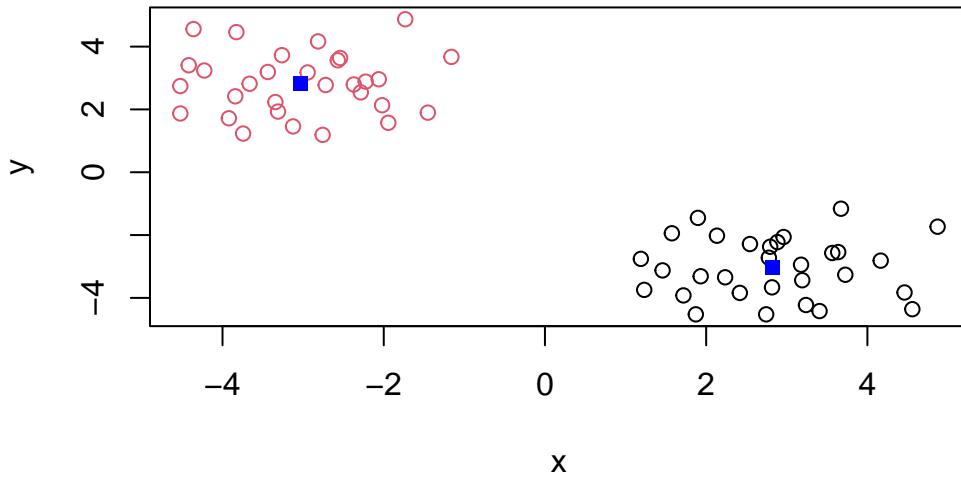
[1] 30 30

Q. What clusters do my data points reside in? What are my cluster results?

k\$cluster

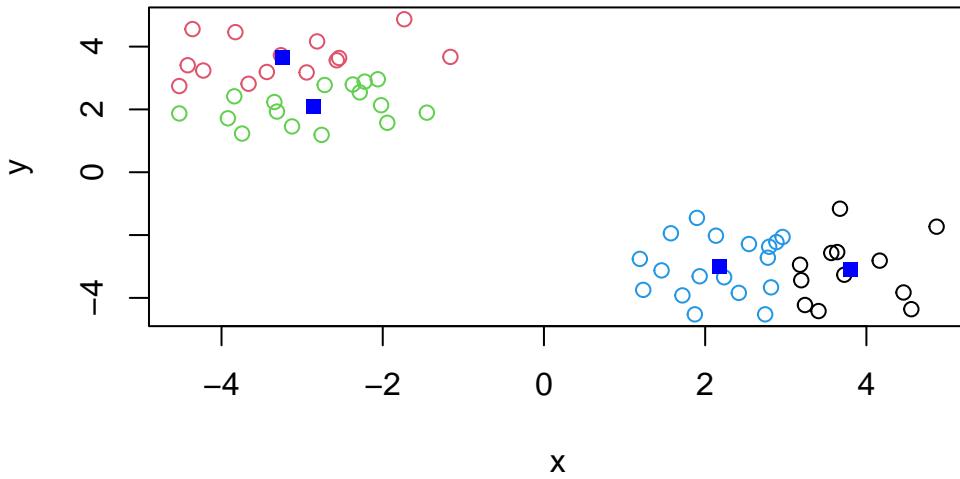
Q. Make a plot of our data colored by cluster assignment - i.e. Make a result figure...

```
plot(x, col=k$cluster)
points(k$centers, col="blue", pch=15)
```



Q. Cluster with k-means into 4 clusters and plot your results as above.

```
k4 <- kmeans(x, centers=4)
plot(x, col=k4$cluster)
points(k4$centers,col="blue", pch=15)
```



Q. Run kmeans with centers (i.e. values of k) equal 1 to 6

```
k$tot.withinss
```

```
[1] 109.6286
```

```
k1 <- kmeans(x, centers=1)$tot.withinss
k2 <- kmeans(x, centers=1)$tot.withinss
k3 <- kmeans(x, centers=1)$tot.withinss
k4 <- kmeans(x, centers=1)$tot.withinss
k5 <- kmeans(x, centers=1)$tot.withinss
k6 <- kmeans(x, centers=1)$tot.withinss

ans <- c(k1,k2,k3,k4,k5,k6)
```

Or use a for loop

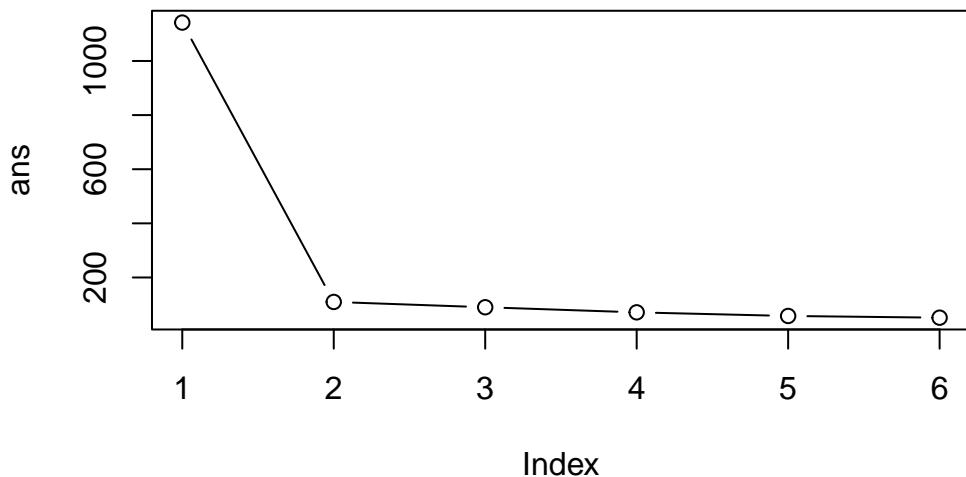
```
ans <- NULL
for(i in 1:6) {
  ans <- c(ans, kmeans(x, centers=i)$tot.withinss)
}
```

Make a “scree-plot’

```
ans
```

```
[1] 1141.87158 109.62862 90.02583 71.17855 57.82773 51.47343
```

```
plot(ans, typ="b")
```



## Hierchiacal Clustering

The main function in “base” R for this is called `hclust()`

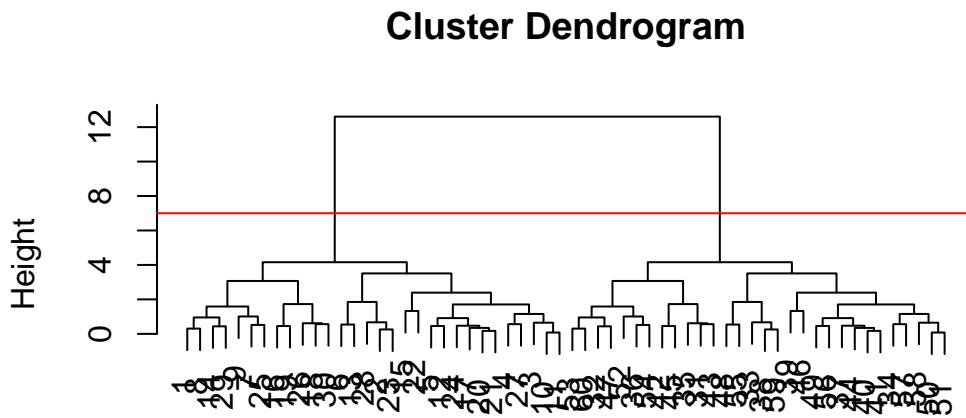
```
d <- dist(x)
hc <- hclust(d)
hc
```

```
Call:
hclust(d = d)
```

```
Cluster method : complete
```

Distance : euclidean  
Number of objects: 60

```
plot(hc)
abline(h=7, col='red')
```

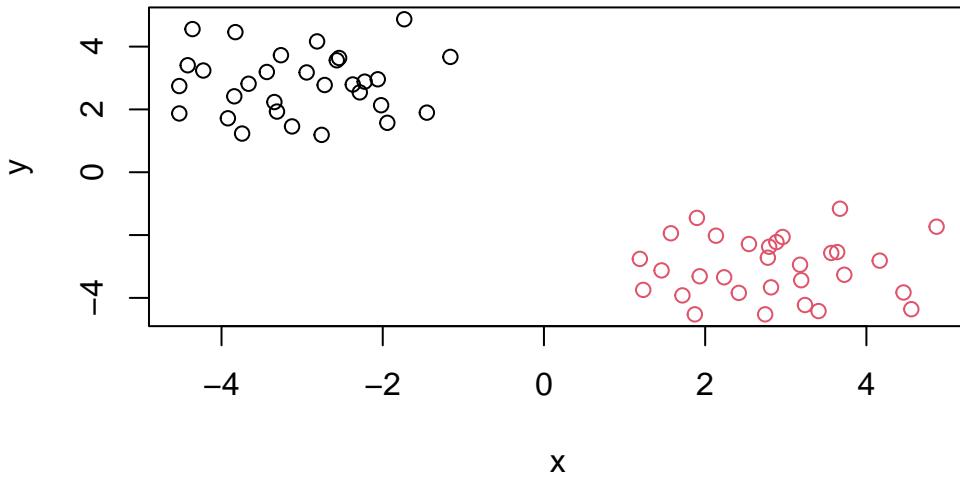


```
d  
hclust (*, "complete")
```

To obtain clusters from our `hclust` result object `hc` we “cut” the tree to yield different sub branches. For this we use the `cutree()` function.

```
grps <- cutree(hc, h=7)  
grps
```

```
plot(x, col=grps)
```



## Principal Component Analysis (PCA)

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this question?

```
dim(x)
```

```
[1] 17 5
```

### Preview the first 6 rows

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66

2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

Fix the amount of rows

```
# Note how the minus indexing works
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Now lets check the dimensions again:

```
dim(x)
```

```
[1] 17 4
```

Side-note: An alternative approach to setting the correct row-names in this case would be to read the data file again and this time set the row.names argument of read.csv() to be the first column (i.e. use argument setting row.names=1), see below:

```
x <- read.csv(url, row.names=1)
head(x)
```

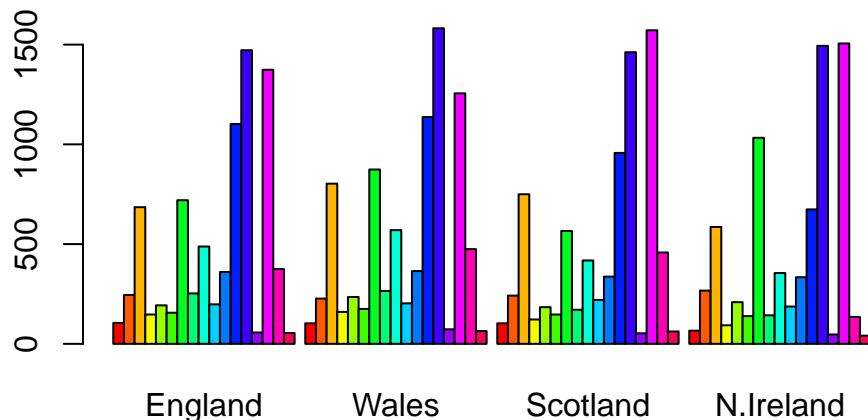
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

If you run `x <- x[, -1]` multiple times, you remove a column every time you run a code chunk. So the best way to run this is reading the data the correct way using `x <- read.csv(url, row.names=1)`

## Spotting major differences and trends

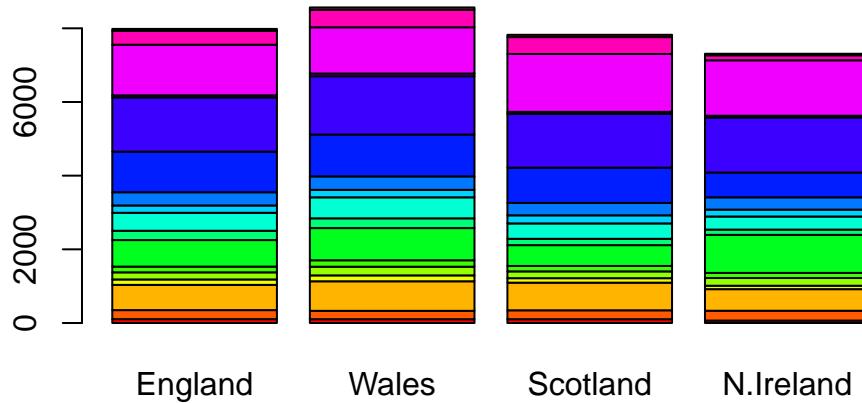
```
# Using base R  
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?

If you change the `beside=T` to `beside=F`

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



## Currently we have wide format

Data organized with one row per Food (17 observations) and multiple columns for Country (4 different measurements across a row):

```
dim(x)
```

```
[1] 17 4
```

To convert this to “long” format we want one row per measurement - maximizes rows ( $17 \times 4 = 68$ ), minimizes columns (with a single Consumption measurement value per Country). We will do tidying with the `pivot_longer()` function from the `tidyverse` package:

```
library(tidyverse)

# Convert data to long format for ggplot with `pivot_longer()`
x_long <- x |>
  tibble::rownames_to_column("Food") |>
  pivot_longer(cols = -Food,
               names_to = "Country",
```

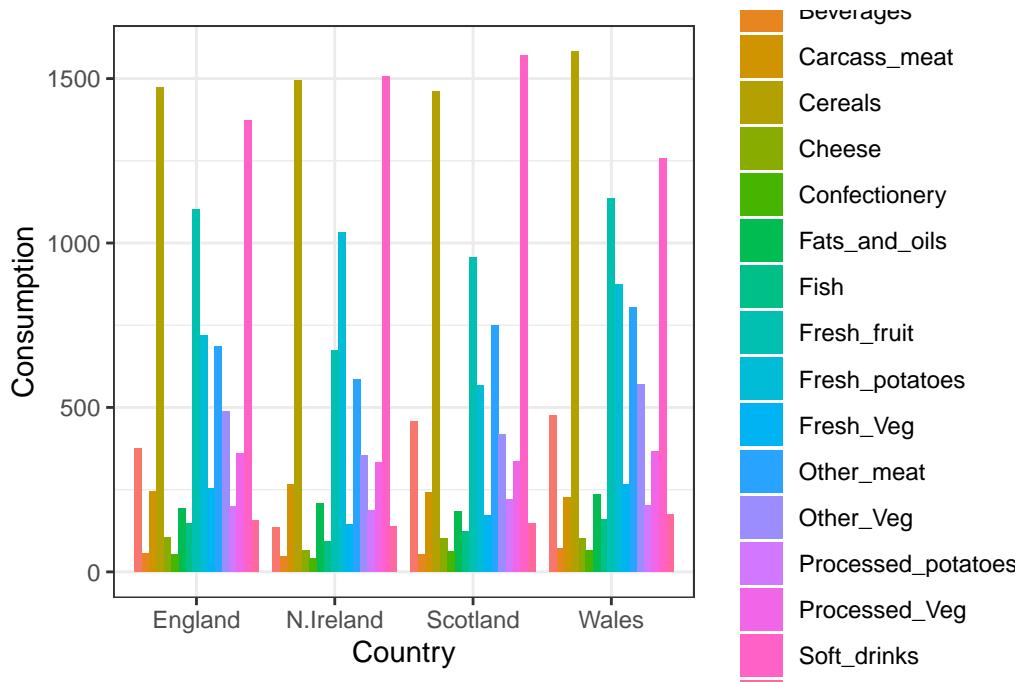
```
    values_to = "Consumption")  
  
dim(x_long)
```

```
[1] 68 3
```

```
head(x_long)
```

```
# A tibble: 6 x 3  
Food          Country  Consumption  
<chr>        <chr>       <int>  
1 "Cheese"     England      105  
2 "Cheese"     Wales        103  
3 "Cheese"     Scotland     103  
4 "Cheese"     N.Ireland     66  
5 "Carcass_meat" England     245  
6 "Carcass_meat" Wales       227
```

```
# Create grouped bar plot  
library(ggplot2)  
  
ggplot(x_long) +  
  aes(x = Country, y = Consumption, fill = Food) +  
  geom_col(position = "dodge") +  
  theme_bw()
```

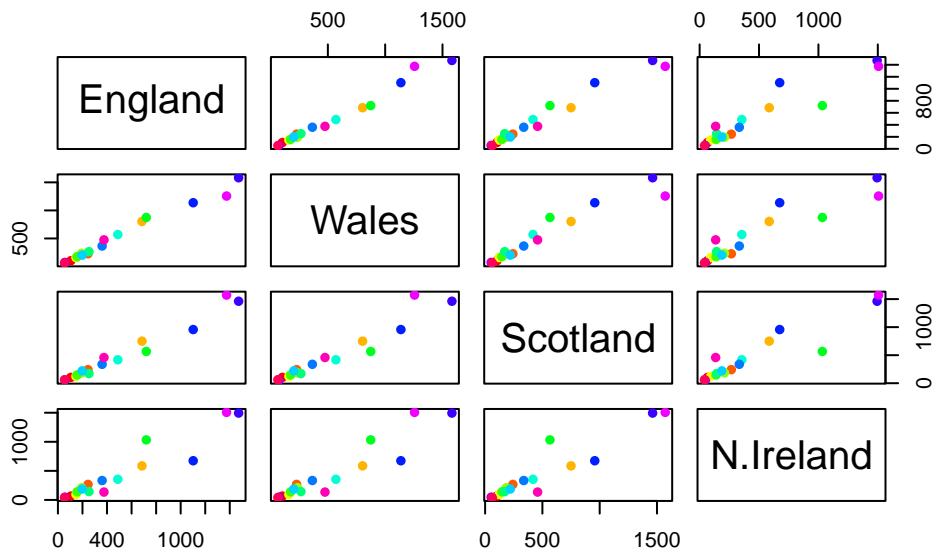


Q4: Changing what optional argument in the above ggplot() code results in a stacked barplot figure?

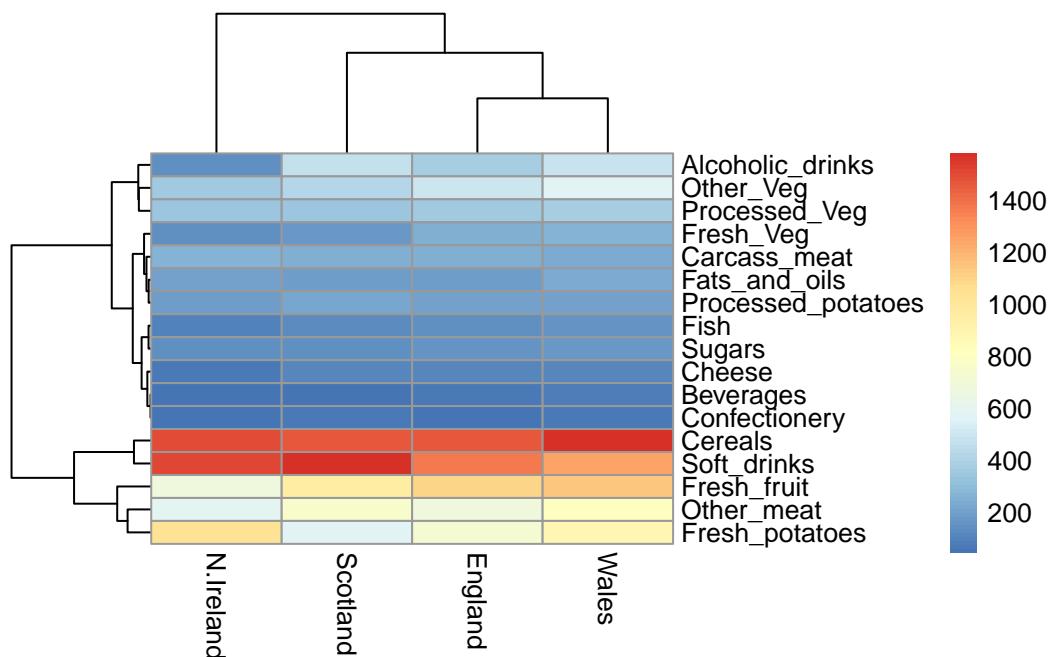
Q5: We can use the pairs() function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

In the y axis of every plot in the first row, England is on the y-axis and Wales on the x-axis, the same pattern goes for all the rest of the rows. The points represent all the different foods. Points on the diagonal mean these foods are consumed similar between countries.

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

It looks like Wales and England are quite similar in their consumption of these foods. it is still quite difficult to tell what is going on in the dataset.

## PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`.

As we want to do PCA on the food data for the different countries we will want the foods in the columns.

```
pca <- prcomp( t(x) )
summary(pca)
```

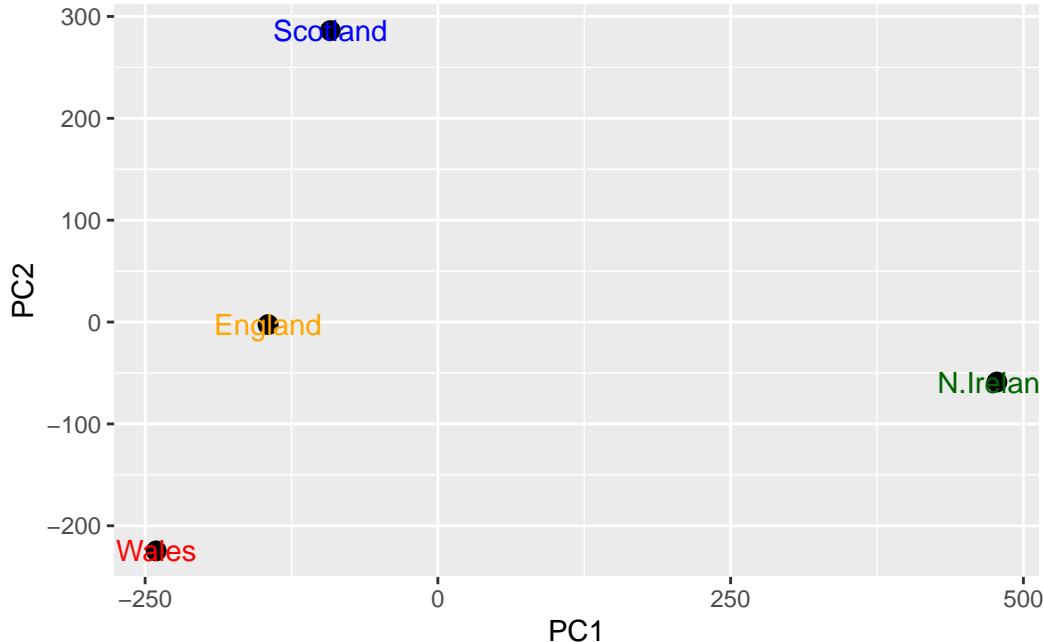
Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Our result object is called `pca` and it has a `$x` component that we will look at first

Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
library(ggplot2)
cols <- c("orange", "red", "blue", "darkgreen")
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(size = 3) +
  geom_text(col=cols)
```



We can use the square of `pca$sdev`, which stands for “standard deviation”, to calculate how much variation in the original data each PC accounts for.

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```

```
## or the second row here...
z <- summary(pca)
z$importance
```

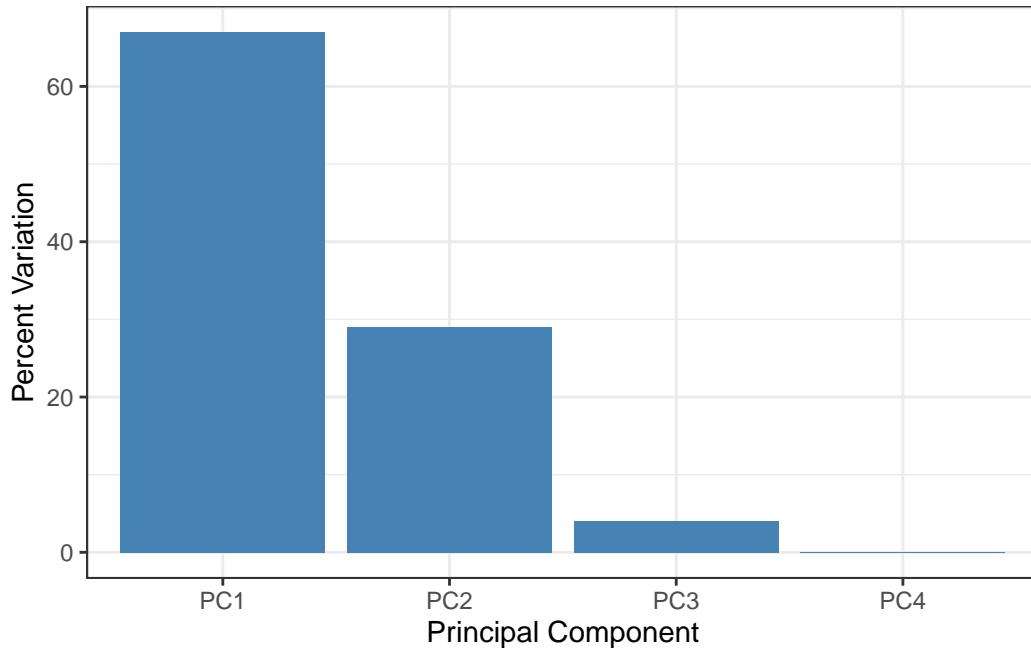
	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.921348e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
# Create scree plot with ggplot
variance_df <- data.frame(
  PC = factor(paste0("PC", 1:length(v)), levels = paste0("PC", 1:length(v))),
  Variance = v
)
```

```

ggplot(variance_df) +
  aes(x = PC, y = Variance) +
  geom_col(fill = "steelblue") +
  xlab("Principal Component") +
  ylab("Percent Variation") +
  theme_bw() +
  theme(axis.text.x = element_text(angle = 90))

```



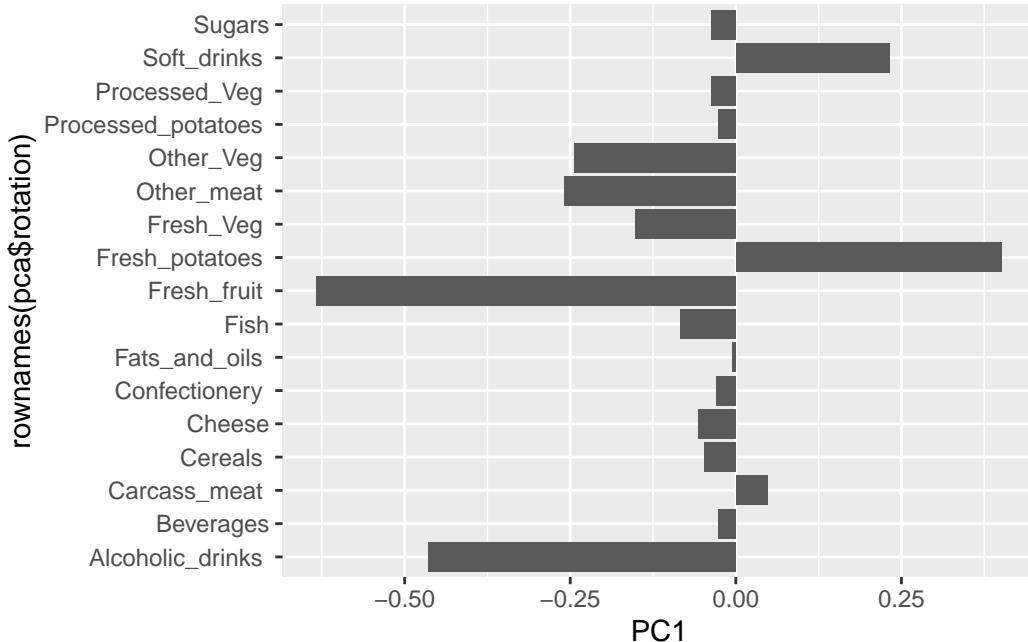
### Digging deeper (variable loadings)

Another major result out of PCA is the so-called “variable loadings” or `$rotation` that tells us how the original variables (foods) contribute to PCs (i.e. our new axis)

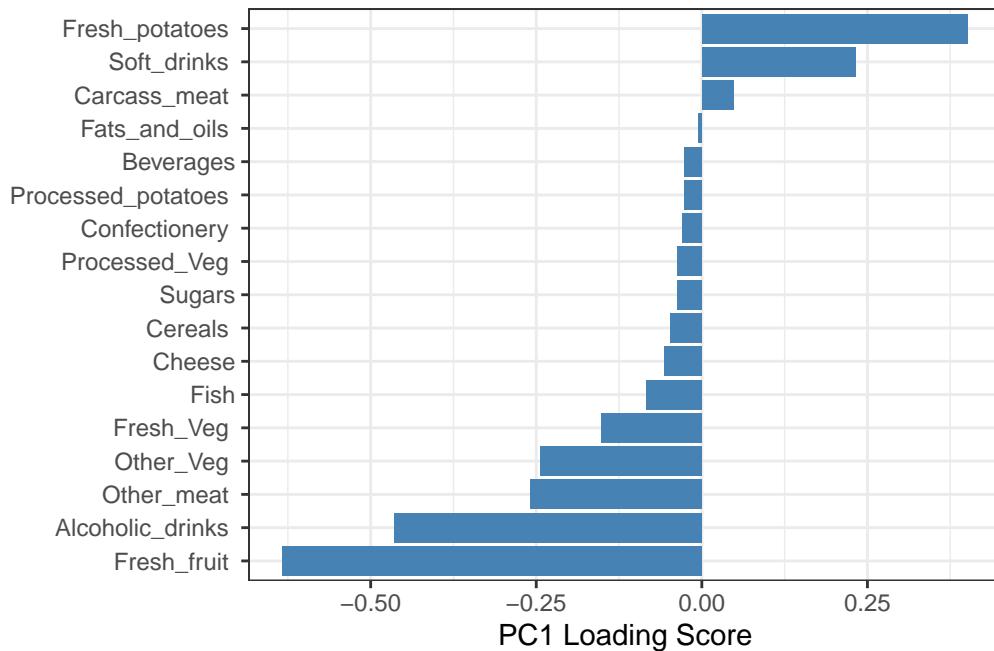
```

ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()

```



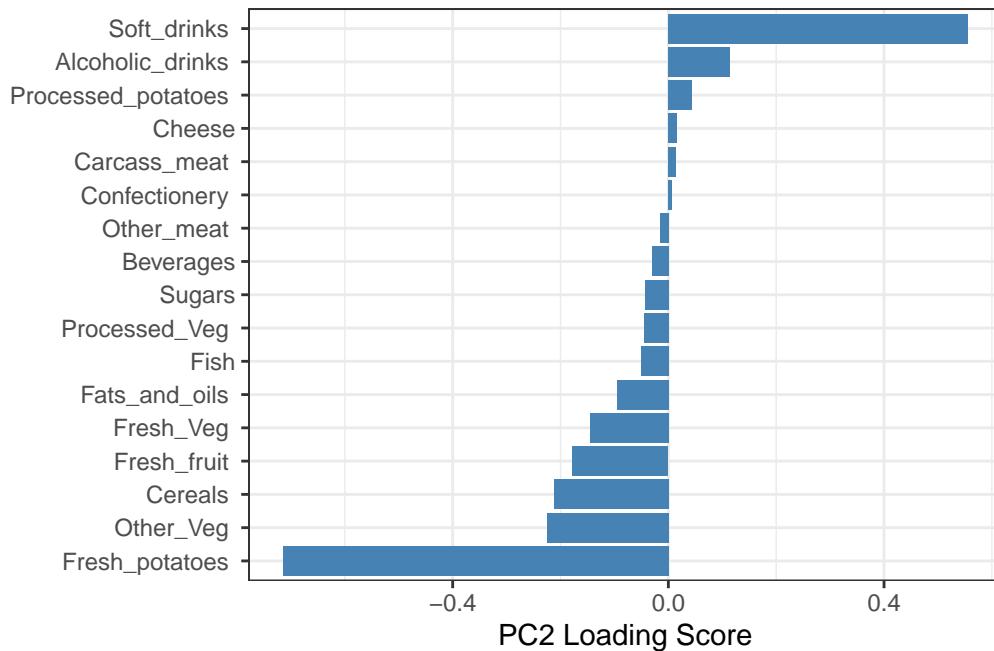
```
## Lets focus on PC1 as it accounts for > 90% of variance
ggplot(pca$rotation) +
  aes(x = PC1,
      y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



Q9: Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

We can see when we plot PC2 there are two outliers, one being for fresh potatoes and the other soft drinks

```
ggplot(pca$rotation) +
  aes(x = PC2,
      y = reorder(rownames(pca$rotation), PC2)) +
  geom_col(fill = "steelblue") +
  xlab("PC2 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



## PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10: How many genes and samples are in this data set?

There are 100 rows, and 10 columns

```
dim(rna.data)
```

```
[1] 100 10
```

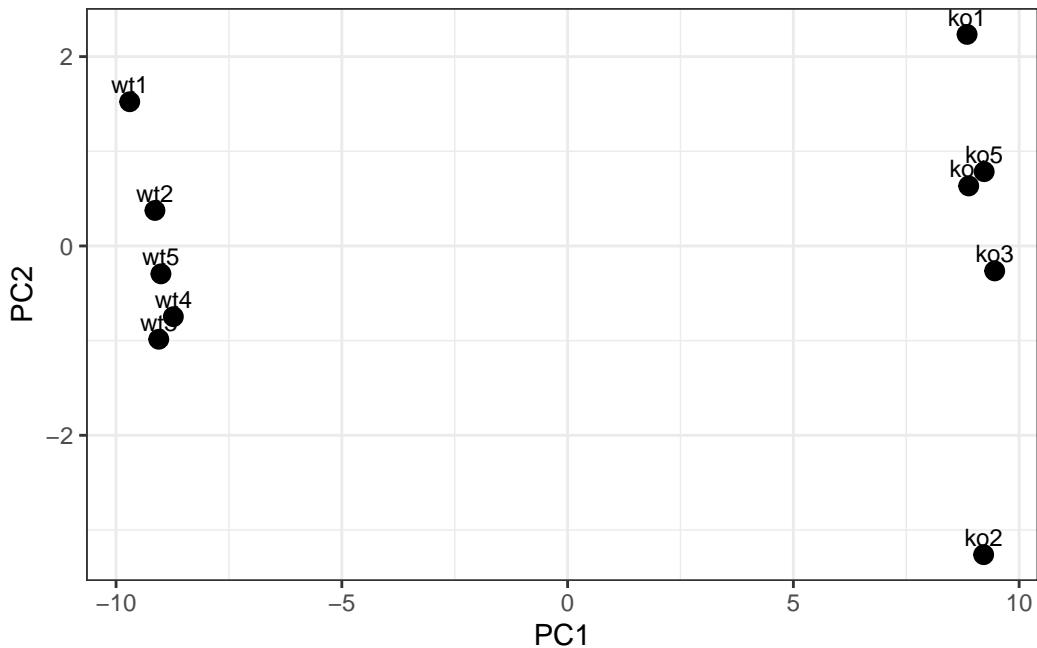
```

## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

# Create data frame for plotting
df <- as.data.frame(pca$x)
df$Sample <- rownames(df)

## Plot with ggplot
ggplot(df) +
  aes(x = PC1, y = PC2, label = Sample) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, size = 3) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()

```



```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642

```

Cumulative Proportion  0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251
                           PC8      PC9      PC10
Standard deviation     0.62065 0.60342 3.3e-15
Proportion of Variance 0.00385 0.00364 0.0e+00
Cumulative Proportion  0.99636 1.00000 1.0e+00

```

A quick scree plot summary of this Proportion of Variance for each PC can be obtained using ggplot:

```

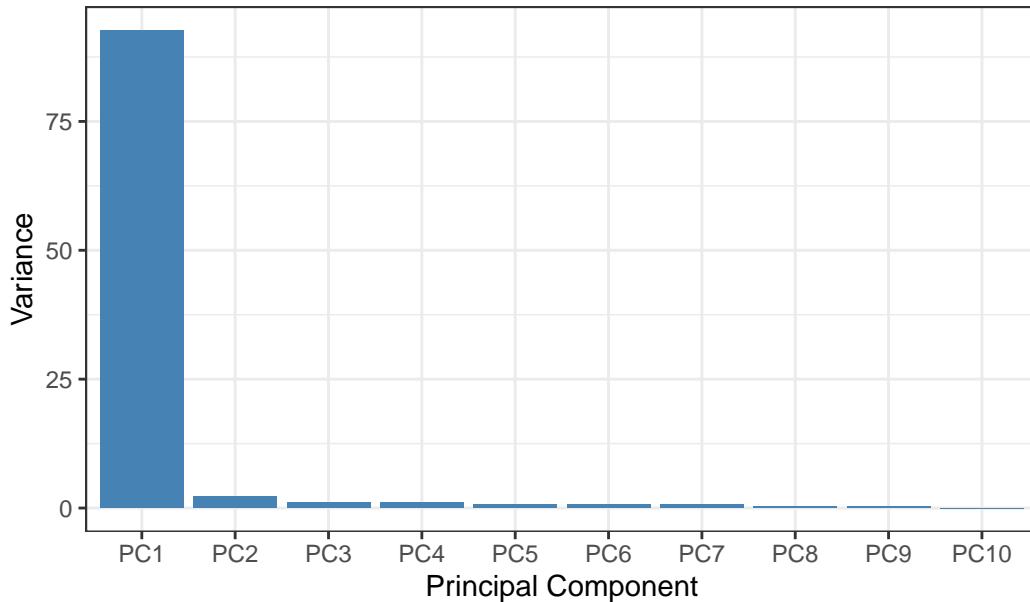
# Calculate variance explained
pca.var <- pca$sdev^2
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)

# Create scree plot data
scree_df <- data.frame(
  PC = factor(paste0("PC", 1:10), levels = paste0("PC", 1:10)),
  Variance = pca.var[1:10]
)

ggplot(scree_df) +
  aes(x = PC, y = Variance) +
  geom_col(fill = "steelblue") +
  ggtitle("Quick scree plot") +
  xlab("Principal Component") +
  ylab("Variance") +
  theme_bw()

```

## Quick scree plot



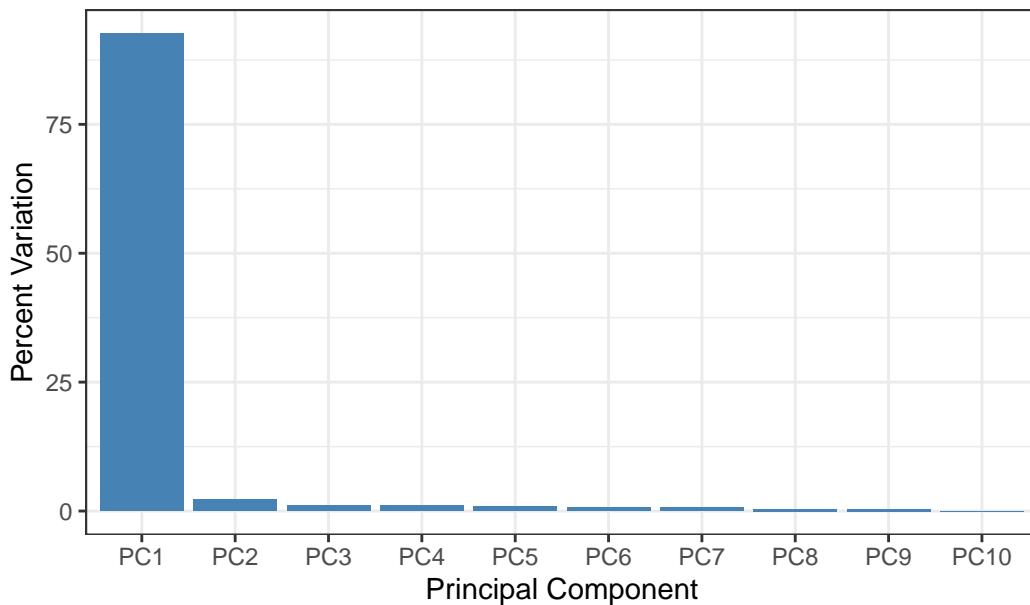
```
## Percent variance is often more informative to look at
pca.var.per
```

```
[1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

```
# Create percent variance scree plot
scree_pct_df <- data.frame(
  PC = factor(paste0("PC", 1:10), levels = paste0("PC", 1:10)),
  PercentVariation = pca.var.per[1:10]
)

ggplot(scree_pct_df) +
  aes(x = PC, y = PercentVariation) +
  geom_col(fill = "steelblue") +
  ggtitle("Scree Plot") +
  xlab("Principal Component") +
  ylab("Percent Variation") +
  theme_bw()
```

## Scree Plot

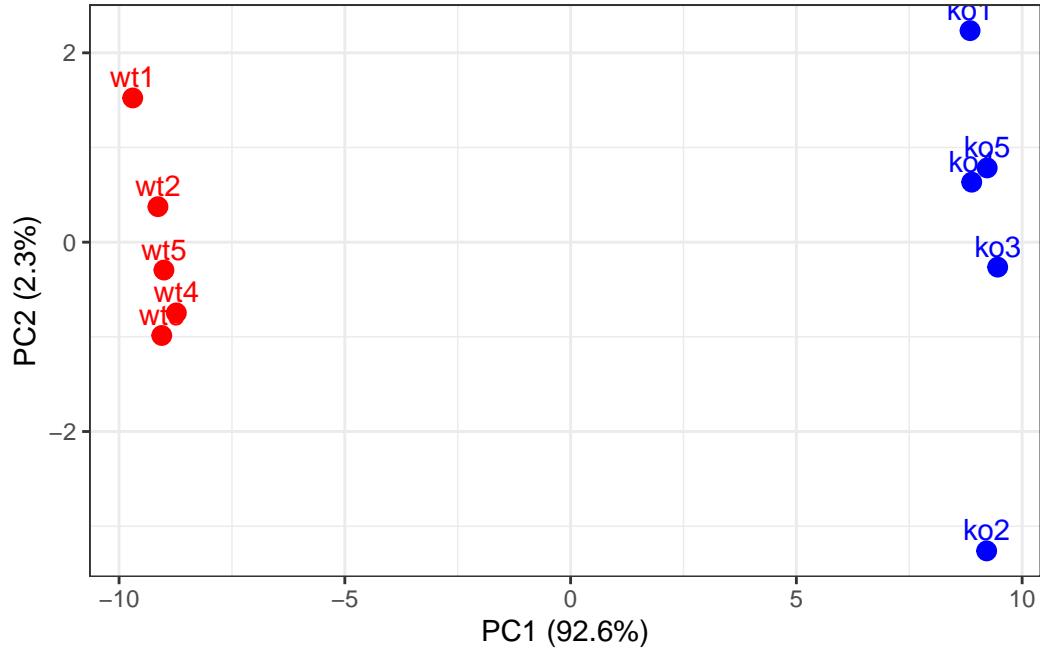


A better looking graph:

```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

# Add condition to data frame
df$condition <- substr(df$Sample, 1, 2)
df$color <- colvec

ggplot(df) +
  aes(x = PC1, y = PC2, color = color, label = Sample) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, hjust = 0.5, show.legend = FALSE) +
  scale_color_identity() +
  xlab(paste0("PC1 (", pca.var.per[1], "%)")) +
  ylab(paste0("PC2 (", pca.var.per[2], "%)")) +
  theme_bw()
```



## Gene Loadings

```

loading_scores <- pca$rotation[,1]

## Find the top 10 measurements (genes) that contribute
## most to PC1 in either direction (+ or -)
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)

## show the names of the top 10 genes
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes

[1] "gene100" "gene66"   "gene45"   "gene68"   "gene98"   "gene60"   "gene21"
[8] "gene56"   "gene10"   "gene90"

```